

Home
Pages Classes Methods
Search
Table of Contents
What's Here
Querying
Comparing
Converting
<div>Show/hide navigation</div>
Parent
Numeric
Methods
#%
#*
#**
#+
#-
#-@
#/
#<
#<=
#<=>
#==
#===
#>
#>=
#abs
#angle
#arg
#ceil
#coerce
#denominator
#divmod
#eql?
#fdiv
#finite?
#floor
#hash
#infinite?
#inspect
#magnitude
#modulo
#nan?
#negative?
#next float
#numerator
#phase
#positive?
#prev float
#quo

[#rationalize](#)
[#round](#)
[#to f](#)
[#to i](#)
[#to int](#)
[#to r](#)
[#to s](#)
[#truncate](#)
[#zero?](#)

class Float

A Float object represents a sometimes-inexact real number using the native architecture's double-precision floating point representation.

Floating point has a different arithmetic and is an inexact number. So you should know its esoteric system. See following:

- docs.oracle.com/cd/E19957-01/806-3568/ncg_goldberg.html
- github.com/rdp/ruby_tutorials_core/wiki/Ruby-Talk-FAQ#-why-are-rubys-floats-imprecise
- en.wikipedia.org/wiki/Floating_point#Accuracy_problems

You can create a Float object explicitly with:

- A [floating-point literal](#).

You can convert certain objects to Floats with:

- Method [Float](#).

What's Here

First, what's elsewhere. Class Float:

- Inherits from [class Numeric](#).

Here, class Float provides methods for:

- [Querying](#)
- [Comparing](#)
- [Converting](#)

Querying

- [finite?](#): Returns whether `self` is finite.
- [hash](#): Returns the integer hash code for `self`.
- [infinite?](#): Returns whether `self` is infinite.
- [nan?](#): Returns whether `self` is a NaN (not-a-number).

Comparing

- `#<`: Returns whether `self` is less than the given value.
- `#<=`: Returns whether `self` is less than or equal to the given value.
- `#<=>`: Returns a number indicating whether `self` is less than, equal to, or greater than the given value.
- `==` (aliased as `===` and [eql?](#)): Returns whether `self` is equal to the given value.
- `#>`: Returns whether `self` is greater than the given value.
- `#>=`: Returns whether `self` is greater than or equal to the given value.

Converting

- `%` (aliased as [modulo](#)): Returns `self` modulo the given value.
- `*`: Returns the product of `self` and the given value.
- `*^`: Returns the value of `self` raised to the power of the given value.
- `+`: Returns the sum of `self` and the given value.
- `-`: Returns the difference of `self` and the given value.
- `/`: Returns the quotient of `self` and the given value.
- [ceil](#): Returns the smallest number greater than or equal to `self`.

- [coerce](#): Returns a 2-element array containing the given value converted to a Float and `self`
- [divmod](#): Returns a 2-element array containing the quotient and remainder results of dividing `self` by the given value.
- [fdiv](#): Returns the Float result of dividing `self` by the given value.
- [floor](#): Returns the greatest number smaller than or equal to `self`.
- [next_float](#): Returns the next-larger representable Float.
- [prev_float](#): Returns the next-smaller representable Float.
- [quo](#): Returns the quotient from dividing `self` by the given value.
- [round](#): Returns `self` rounded to the nearest value, to a given precision.
- [to_i](#) (aliased as [to_int](#)): Returns `self` truncated to an [Integer](#).
- [to_s](#) (aliased as [inspect](#)): Returns a string containing the place-value representation of `self` in the given radix.
- [truncate](#): Returns `self` truncated to a given precision.

Constants

DIG

The minimum number of significant decimal digits in a double-precision floating point.

Usually defaults to 15.

EPSILON

The difference between 1 and the smallest double-precision floating point number greater than 1.

Usually defaults to 2.2204460492503131e-16.

INFINITY

An expression representing positive infinity.

MANT_DIG

The number of base digits for the `double` data type.

Usually defaults to 53.

MAX

The largest possible integer in a double-precision floating point number.

Usually defaults to 1.7976931348623157e+308.

MAX_10_EXP

The largest positive exponent in a double-precision floating point where 10 raised to this power minus 1.

Usually defaults to 308.

MAX_EXP

The largest possible exponent value in a double-precision floating point.

Usually defaults to 1024.

MIN

The smallest positive normalized number in a double-precision floating point.

Usually defaults to 2.2250738585072014e-308.

If the platform supports denormalized numbers, there are numbers between zero and [Float::MIN](#). `0.0.next_float` returns the smallest positive floating point number including denormalized numbers.

MIN_10_EXP

The smallest negative exponent in a double-precision floating point where 10 raised to this power minus 1.

Usually defaults to -307.

MIN_EXP

The smallest possible exponent value in a double-precision floating point.

Usually defaults to -1021.

NAN

An expression representing a value which is “not a number”.

RADIX

The base of the floating point, or number of unique digits used to represent the number.

Usually defaults to 2 on most systems, which would represent a base-10 decimal.

Public Instance Methods

self % other → float

Returns `self` modulo `other` as a float.

For float `f` and real number `r`, these expressions are equivalent:

```
f % r
f-r*(f/r).floor
f.divmod(r)[1]
```

See [Numeric#divmod](#).

Examples:

```
10.0 % 2          # => 0.0
10.0 % 3          # => 1.0
10.0 % 4          # => 2.0

10.0 % -2         # => 0.0
10.0 % -3         # => -2.0
10.0 % -4         # => -2.0

10.0 % 4.0        # => 2.0
10.0 % Rational(4, 1) # => 2.0
```

Also aliased as: [modulo](#)

self * other → numeric

Returns a new Float which is the product of `self` and `other`:

```
f = 3.14
f * 2          # => 6.28
f * 2.0        # => 6.28
f * Rational(1, 2) # => 1.57
f * Complex(2, 0)  # => (6.28+0.0i)
```

self ** other → numeric

Raises `self` to the power of `other`:

```
f = 3.14
f ** 2          # => 9.8596
f ** -2         # => 0.1014239928597509
f ** 2.1        # => 11.054834900588839
f ** Rational(2, 1) # => 9.8596
f ** Complex(2, 0)  # => (9.8596+0i)
```

self + other → numeric

Returns a new Float which is the sum of `self` and `other`:

```
f = 3.14
f + 1           # => 4.1400000000000001
f + 1.0         # => 4.1400000000000001
f + Rational(1, 1) # => 4.1400000000000001
f + Complex(1, 0)  # => (4.1400000000000001+0i)
```

self - other → numeric

Returns a new Float which is the difference of `self` and `other`:

```
f = 3.14
f - 1           # => 2.14
f - 1.0         # => 2.14
f - Rational(1, 1) # => 2.14
f - Complex(1, 0)  # => (2.14+0i)
```

-float → float

Returns `self`, negated.

self / other → numeric

Returns a new Float which is the result of dividing `self` by `other`:

```
f = 3.14
f / 2           # => 1.57
f / 2.0         # => 1.57
f / Rational(2, 1) # => 1.57
f / Complex(2, 0)  # => (1.57+0.0i)
```

self < other → true or false

Returns `true` if `self` is numerically less than `other`:

```
2.0 < 3           # => true
2.0 < 3.0         # => true
2.0 < Rational(3, 1) # => true
2.0 < 2.0         # => false
```

`Float::NAN < Float::NAN` returns an implementation-dependent value.

self <= other → true or false

Returns `true` if `self` is numerically less than or equal to `other` :

```
2.0 <= 3           # => true
2.0 <= 3.0         # => true
2.0 <= Rational(3, 1) # => true
2.0 <= 2.0         # => true
2.0 <= 1.0         # => false
```

`Float::NAN <= Float::NAN` returns an implementation-dependent value.

self <=> other → -1, 0, +1, or nil

Returns a value that depends on the numeric relation between `self` and `other` :

- -1, if `self` is less than `other` .
- 0, if `self` is equal to `other` .
- 1, if `self` is greater than `other` .
- `nil` , if the two values are incommensurate.

Examples:

```
2.0 <=> 2           # => 0
2.0 <=> 2.0         # => 0
2.0 <=> Rational(2, 1) # => 0
2.0 <=> Complex(2, 0) # => 0
2.0 <=> 1.9         # => 1
2.0 <=> 2.1         # => -1
2.0 <=> 'foo'       # => nil
```

This is the basis for the tests in the [Comparable](#) module.

`Float::NAN <=> Float::NAN` returns an implementation-dependent value.

self == other → true or false

Returns `true` if `other` has the same value as `self`, `false` otherwise:

```
2.0 == 2           # => true
2.0 == 2.0         # => true
2.0 == Rational(2, 1) # => true
2.0 == Complex(2, 0) # => true
```

`Float::NAN == Float::NAN` returns an implementation-dependent value.

Related: [Float#eq?l?](#) (requires `other` to be a `Float`).

Also aliased as: [===](#)

==(p1)

Returns `true` if `other` has the same value as `self`, `false` otherwise:

```
2.0 == 2           # => true
2.0 == 2.0         # => true
2.0 == Rational(2, 1) # => true
2.0 == Complex(2, 0) # => true
```

`Float::NAN == Float::NAN` returns an implementation-dependent value.

Related: [Float#eql?](#) (requires `other` to be a `Float`).

Alias for: [==](#)

self > other → true or false

Returns `true` if `self` is numerically greater than `other`:

```
2.0 > 1           # => true
2.0 > 1.0         # => true
2.0 > Rational(1, 2) # => true
2.0 > 2.0         # => false
```

`Float::NAN > Float::NAN` returns an implementation-dependent value.

self >= other → true or false

Returns `true` if `self` is numerically greater than or equal to `other`:

```
2.0 >= 1           # => true
2.0 >= 1.0         # => true
2.0 >= Rational(1, 2) # => true
2.0 >= 2.0         # => true
2.0 >= 2.1         # => false
```

`Float::NAN >= Float::NAN` returns an implementation-dependent value.

abs → float

Returns the absolute value of `self`:

```
(-34.56).abs # => 34.56
-34.56.abs  # => 34.56
34.56.abs   # => 34.56
```

angle → 0 or float

Returns 0 if the value is positive, pi otherwise.

Alias for: [arg](#)

arg → 0 or float

Returns 0 if the value is positive, pi otherwise.

Also aliased as: [angle](#), [phase](#)

ceil(ndigits = 0) → float or integer

Returns the smallest number greater than or equal to `self` with a precision of `ndigits` decimal digits.

When `ndigits` is positive, returns a float with `ndigits` digits after the decimal point (as available):

```
f = 12345.6789
f.ceil(1) # => 12345.7
f.ceil(3) # => 12345.679
f = -12345.6789
f.ceil(1) # => -12345.6
f.ceil(3) # => -12345.678
```

When `ndigits` is non-positive, returns an integer with at least `ndigits.abs` trailing zeros:

```
f = 12345.6789
f.ceil(0) # => 12346
f.ceil(-3) # => 13000
f = -12345.6789
f.ceil(0) # => -12345
f.ceil(-3) # => -12000
```

Note that the limited precision of floating-point arithmetic may lead to surprising results:

```
(2.1 / 0.7).ceil #=> 4 (!)
```

Related: [Float#floor](#).

coerce(other) → array

Returns a 2-element array containing `other` converted to a Float and `self`:

```
f = 3.14           # => 3.14
f.coerce(2)       # => [2.0, 3.14]
f.coerce(2.0)     # => [2.0, 3.14]
f.coerce(Rational(1, 2)) # => [0.5, 3.14]
f.coerce(Complex(1, 0)) # => [1.0, 3.14]
```

Raises an exception if a type conversion fails.

denominator → **integer**

Returns the denominator (always positive). The result is machine dependent.

See also [Float#numerator](#).

divmod(other) → **array**

Returns a 2-element array [q, r], where

```
q = (self/other).floor    # Quotient
r = self % other          # Remainder
```

Examples:

```
11.0.divmod(4)           # => [2, 3.0]
11.0.divmod(-4)          # => [-3, -1.0]
-11.0.divmod(4)          # => [-3, 1.0]
-11.0.divmod(-4)         # => [2, -3.0]

12.0.divmod(4)           # => [3, 0.0]
12.0.divmod(-4)          # => [-3, 0.0]
-12.0.divmod(4)          # => [-3, -0.0]
-12.0.divmod(-4)         # => [3, -0.0]

13.0.divmod(4.0)         # => [3, 1.0]
13.0.divmod(Rational(4, 1)) # => [3, 1.0]
```

eq?(other) → **true or false**

Returns **true** if **other** is a Float with the same value as **self**, **false** otherwise:

```
2.0.eq?(2.0)           # => true
2.0.eq?(1.0)           # => false
2.0.eq?(1)             # => false
2.0.eq?(Rational(2, 1)) # => false
2.0.eq?(Complex(2, 0)) # => false
```

`Float::NAN.eq?(Float::NAN)` returns an implementation-dependent value.

Related: [Float#==](#) (performs type conversions).

fdiv(p1)

Returns the quotient from dividing `self` by `other`:

```
f = 3.14
f.quo(2)           # => 1.57
f.quo(-2)          # => -1.57
f.quo(Rational(2, 1)) # => 1.57
f.quo(Complex(2, 0)) # => (1.57+0.0i)
```

Alias for: [quo](#)

finite? → true or false

Returns `true` if `self` is not `Infinity`, `-Infinity`, or `NaN`, `false` otherwise:

```
f = 2.0           # => 2.0
f.finite?        # => true
f = 1.0/0.0       # => Infinity
f.finite?        # => false
f = -1.0/0.0      # => -Infinity
f.finite?        # => false
f = 0.0/0.0       # => NaN
f.finite?        # => false
```

floor(ndigits = 0) → float or integer

Returns the largest number less than or equal to `self` with a precision of `ndigits` decimal digits.

When `ndigits` is positive, returns a float with `ndigits` digits after the decimal point (as available):

```
f = 12345.6789
f.floor(1) # => 12345.6
f.floor(3) # => 12345.678
f = -12345.6789
f.floor(1) # => -12345.7
f.floor(3) # => -12345.679
```

When `ndigits` is non-positive, returns an integer with at least `ndigits.abs` trailing zeros:

```
f = 12345.6789
f.floor(0) # => 12345
f.floor(-3) # => 12000
f = -12345.6789
f.floor(0) # => -12346
f.floor(-3) # => -13000
```

Note that the limited precision of floating-point arithmetic may lead to surprising results:

```
(0.3 / 0.1).floor #=> 2 (!)
```

Related: [Float#ceil](#).

hash → integer

Returns the integer hash value for `self`.

See also [Object#hash](#).

infinite? → -1, 1, or nil

Returns:

- 1, if `self` is `Infinity`.
- -1 if `self` is `-Infinity`.
- `nil`, otherwise.

Examples:

```
f = 1.0/0.0 # => Infinity
f.infinite? # => 1
f = -1.0/0.0 # => -Infinity
f.infinite? # => -1
f = 1.0 # => 1.0
f.infinite? # => nil
f = 0.0/0.0 # => NaN
f.infinite? # => nil
```

inspect()

Returns a string containing a representation of `self`; depending of the value of `self`, the string representation may contain:

- A fixed-point number.
- A number in “scientific notation” (containing an exponent).
- ‘Infinity’.
- ‘-Infinity’.
- ‘NaN’ (indicating not-a-number).

```
3.14.to_s # => "3.14" (10.1**50).to_s # => "1.644631821843879e+50"
(10.1**500).to_s # => "Infinity" (-10.1**500).to_s # => "-Infinity" (0.0/0.0).to_s # =>
```

“NaN”

Alias for: [to_s](#)

magnitude()

modulo(p1)

Returns `self` modulo `other` as a float.

For float `f` and real number `r`, these expressions are equivalent:

```
f % r
f-r*(f/r).floor
f.divmod(r)[1]
```

See [Numeric#divmod](#).

Examples:

```
10.0 % 2          # => 0.0
10.0 % 3          # => 1.0
10.0 % 4          # => 2.0

10.0 % -2         # => 0.0
10.0 % -3         # => -2.0
10.0 % -4         # => -2.0

10.0 % 4.0        # => 2.0
10.0 % Rational(4, 1) # => 2.0
```

Alias for: [%](#)

nan? → true or false

Returns `true` if `self` is a NaN, `false` otherwise.

```
f = -1.0          #=> -1.0
f.nan?            #=> false
f = 0.0/0.0       #=> NaN
f.nan?            #=> true
```

negative? → true or false

Returns `true` if `self` is less than 0, `false` otherwise.

next_float → float

Returns the next-larger representable Float.

These examples show the internally stored values (64-bit hexadecimal) for each Float `f` and for the corresponding `f.next_float`:

```
f = 0.0      # 0x0000000000000000
f.next_float # 0x0000000000000001

f = 0.01     # 0x3f847ae147ae147b
f.next_float # 0x3f847ae147ae147c
```

In the remaining examples here, the output is shown in the usual way (result `to_s`):

```
0.01.next_float # => 0.010000000000000002
1.0.next_float  # => 1.0000000000000002
100.0.next_float # => 100.00000000000001

f = 0.01
(0..3).each_with_index {|i| printf "%2d %-20a %s\n", i, f, f.to_s; f = f.next_float}
```

Output:

```
0 0x1.47ae147ae147bp-7 0.01
1 0x1.47ae147ae147cp-7 0.010000000000000002
2 0x1.47ae147ae147dp-7 0.010000000000000004
3 0x1.47ae147ae147ep-7 0.010000000000000005

f = 0.0; 100.times { f += 0.1 }
f                                # => 9.999999999999998      # should be 10.0 in the
10-f                            # => 1.9539925233402755e-14 # the floating point error
10.0.next_float-10              # => 1.7763568394002505e-15 # 1 ulp (unit in the last place)
(10-f)/(10.0.next_float-10)     # => 11.0                  # the error is 11 ulps
(10-f)/(10*Float::EPSILON)      # => 8.8                   # approximation of the error
"%a" % 10                       # => "0x1.4p+3"
"%a" % f                        # => "0x1.3ffffffffff5p+3" # the last hex digit is 5
```

Related: [Float#prev_float](#)

numerator → integer

Returns the numerator. The result is machine dependent.

```
n = 0.3.numerator #=> 5404319552844595
d = 0.3.denominator #=> 18014398509481984
n.fdiv(d)          #=> 0.3
```

See also [Float#denominator](#).

phase → 0 or float

Returns 0 if the value is positive, pi otherwise.

Alias for: [arg](#)

positive? → true or false

Returns true if self is greater than 0, false otherwise.

prev_float → float

Returns the next-smaller representable Float.

These examples show the internally stored values (64-bit hexadecimal) for each Float `f` and for the corresponding `f.prev_float`:

```
f = 5e-324      # 0x0000000000000001
f.prev_float   # 0x0000000000000000

f = 0.01        # 0x3f847ae147ae147b
f.prev_float    # 0x3f847ae147ae147a
```

In the remaining examples here, the output is shown in the usual way (result `to_s`):

```
0.01.prev_float   # => 0.009999999999999998
1.0.prev_float    # => 0.9999999999999999
100.0.prev_float  # => 99.99999999999999

f = 0.01
(0..3).each_with_index {|i| printf "%2d %-20a %s\n", i, f, f.to_s; f = f.prev_float}
```

Output:

```
0 0x1.47ae147ae147bp-7 0.01
1 0x1.47ae147ae147ap-7 0.009999999999999998
2 0x1.47ae147ae1479p-7 0.009999999999999997
3 0x1.47ae147ae1478p-7 0.009999999999999995
```

Related: [Float#next_float](#).

quo(other) → numeric

Returns the quotient from dividing self by other:

```
f = 3.14
f.quo(2)      # => 1.57
f.quo(-2)     # => -1.57
```



```
f.quo(Rational(2, 1)) # => 1.57
f.quo(Complex(2, 0)) # => (1.57+0.0i)
```

Also aliased as: [fdiv](#)

rationalize([eps]) → rational

Returns a simpler approximation of the value ($\text{flt} - |\text{eps}| \leq \text{result} \leq \text{flt} + |\text{eps}|$). If the optional argument `eps` is not given, it will be chosen automatically.

```
0.3.rationalize      #=> (3/10)
1.333.rationalize    #=> (1333/1000)
1.333.rationalize(0.01) #=> (4/3)
```

See also [Float#to_r](#).

round(ndigits = 0, half: :up]) → integer or float

Returns `self` rounded to the nearest value with a precision of `ndigits` decimal digits.

When `ndigits` is non-negative, returns a float with `ndigits` after the decimal point (as available):

```
f = 12345.6789
f.round(1) # => 12345.7
f.round(3) # => 12345.679
f = -12345.6789
f.round(1) # => -12345.7
f.round(3) # => -12345.679
```

When `ndigits` is negative, returns an integer with at least `ndigits.abs` trailing zeros:

```
f = 12345.6789
f.round(0) # => 12346
f.round(-3) # => 12000
f = -12345.6789
f.round(0) # => -12346
f.round(-3) # => -12000
```

If keyword argument `half` is given, and `self` is equidistant from the two candidate values, the rounding is according to the given `half` value:

- `:up` or `nil`: round away from zero:

```
2.5.round(half: :up) # => 3
3.5.round(half: :up) # => 4
```

```
(-2.5).round(half: :up) # => -3
```

- **:down** : round toward zero:

```
2.5.round(half: :down) # => 2
3.5.round(half: :down) # => 3
(-2.5).round(half: :down) # => -2
```

- **:even** : round toward the candidate whose last nonzero digit is even:

```
2.5.round(half: :even) # => 2
3.5.round(half: :even) # => 4
(-2.5).round(half: :even) # => -2
```

Raises an exception if the value for `half` is invalid.

Related: [Float#truncate](#).

to_f → self

Returns `self` (which is already a Float).

to_i → integer

Returns `self` truncated to an [Integer](#).

```
1.2.to_i # => 1
(-1.2).to_i # => -1
```

Note that the limited precision of floating-point arithmetic may lead to surprising results:

```
(0.3 / 0.1).to_i # => 2 (!)
```

Also aliased as: [to_int](#)

to_int()

Returns `self` truncated to an [Integer](#).

```
1.2.to_i # => 1
(-1.2).to_i # => -1
```

Note that the limited precision of floating-point arithmetic may lead to surprising results:

```
(0.3 / 0.1).to_i # => 2 (!)
```

Alias for: [to_i](#)

to_r → rational

Returns the value as a rational.

```
2.0.to_r      #=> (2/1)
2.5.to_r      #=> (5/2)
-0.75.to_r    #=> (-3/4)
0.0.to_r      #=> (0/1)
0.3.to_r      #=> (5404319552844595/18014398509481984)
```

NOTE: 0.3.to_r isn't the same as "0.3".to_r. The latter is equivalent to "3/10".to_r, but the former isn't so.

```
0.3.to_r      == 3/10r  #=> false
"0.3".to_r    == 3/10r  #=> true
```

See also [Float#rationalize](#).

to_s → string

Returns a string containing a representation of `self`; depending of the value of `self`, the string representation may contain:

- A fixed-point number.
- A number in “scientific notation” (containing an exponent).
- ‘Infinity’.
- ‘-Infinity’.
- ‘NaN’ (indicating not-a-number).

```
3.14.to_s # => "3.14" (10.1**50).to_s # => "1.644631821843879e+50"
(10.1**500).to_s # => "Infinity" (-10.1**500).to_s # => "-Infinity" (0.0/0.0).to_s # =>
"NaN"
```

Also aliased as: [inspect](#)

truncate(ndigits = 0) → float or integer

Returns `self` truncated (toward zero) to a precision of `ndigits` decimal digits.

When `ndigits` is positive, returns a float with `ndigits` digits after the decimal point (as available):

```
f = 12345.6789
f.truncate(1) # => 12345.6
f.truncate(3) # => 12345.678
f = -12345.6789
f.truncate(1) # => -12345.6
f.truncate(3) # => -12345.678
```

When `ndigits` is negative, returns an integer with at least `ndigits.abs` trailing zeros:

```
f = 12345.6789
f.truncate(0) # => 12345
f.truncate(-3) # => 12000
f = -12345.6789
f.truncate(0) # => -12345
f.truncate(-3) # => -12000
```

Note that the limited precision of floating-point arithmetic may lead to surprising results:

```
(0.3 / 0.1).truncate #=> 2 (!)
```

Related: [Float#round](#).

zero? → true or false

Returns `true` if `self` is 0.0, `false` otherwise.

[Validate](#)

Generated by [RDoc](#) 6.4.0.

Based on [Darkfish](#) by [Michael Granger](#).

[Ruby-doc.org](#) is provided by [James Britt](#) and [Neurogami](#).

[Hack your world. Feed your head. Live curious.](#)